

Edge Computing Lab

Class: TY-AIEC

School of Computing, MIT Art Design Technology University

Academic Year: 2024-25

Experiment No. 8

Name: Viresh Kamlapure

Class: TY AIEC

Enrollment No: MITU23BTCSD082

Roll No: D2233082

Introduction

The "magic wand" project that can recognize gestures using an accelerometer and an ML classification model on Edge Devices

Objective: Build a project to detect the accelerometer values and convert them into gestures

Tasks:

- Generate the dataset for Accelerometer Motion (Up-Down, Left-Right)
- Configure BLE Sense / Mobile for Edge Impulse
- Building and Training a Model
- Deploy on Nano BLE Sense / Mobile Phone

Introduction

Edge Impulse is a development platform for machine learning on edge devices, targeted at developers who want to create intelligent device solutions. The " Accelerometer Motion "sensor reading equivalent in Edge Impulse would typically involve creating a simple machine learning model that can run on an edge device, like classifying sensor data or recognizing a basic pattern.

Materials Required

- Nano BLE Sense Board

Theory

GPIO (General Purpose Input/Output) pins on the Raspberry Pi are used for interfacing with other electronic components. BCM numbering refers to the pin numbers in the Broadcom SOC channel, which is a more consistent way to refer to the GPIO pins across different versions of the

Here's a high-level overview of steps you'd follow to create a "Hello World" project on Edge Impulse:

Steps to Configure the Edge Impulse:

1. Create an Account and New Project:
 - Sign up for an Edge Impulse account.

- Create a new project from the dashboard.
2. Connect a Device:
 - You can use a supported development board or your smartphone as a sensor device.
 - Follow the instructions to connect your device to your Edge Impulse project.
 3. Collect Data:
 - Use the Edge Impulse mobile app or the Web interface to collect data from the onboard sensors.
 - For a "Hello World" project, you could collect accelerometer data, for instance.
 4. Create an Impulse:
 - Go to the 'Create impulse' page.
 - Add a processing block (e.g., time-series data) and a learning block (e.g., classification).
 - Save the impulse, which defines the machine learning pipeline.
 5. Design a Neural Network:
 - Navigate to the 'NN Classifier' under the 'Learning blocks'.
 - Design a simple neural network. Edge Impulse provides a default architecture that works well for most basic tasks.
 6. Train the Model:
 - Click on the 'Start training' button to train your machine learning model with the collected data.
 7. Test the Model:
 - Once the model is trained, you can test its performance with new data in the 'Model Testing' tab.
 8. Deploy the Model:

9. Run Inference:

10. Monitor:

- Paste your Edge Impulse project's Results:

COM11

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 124 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 124 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 124 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 124 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Predictions (DSP: 122 ms., Classification: 0 ms., Anomaly: 0 ms.): ideal [0, 5, 0, 0, 0, 0, 0, 0,]

Autoscroll Show timestamp Newline 9600 baud Clear output

2) Feature extraction - Image

The screenshot shows the EDGE IMPULSE web interface. On the left is a sidebar with navigation links: Dashboard, Devices, Data acquisition, Experiments, EON Tuner, and Impulse design. The main area is titled 'Dataset' and shows 'DATA COLLECTED 5m 37s' and 'TRAIN / TEST SPLIT 76% / 24%'. Below this is a table of training samples.

SAMPLE NAME	LABEL	ADDED	LENGTH
up-down.5me3umm5	up-down	Mar 17 2025, 12...	4s
up-down.5me3u8p6	up-down	Mar 17 2025, 12...	4s
up-down.5me3tm4e	up-down	Mar 17 2025, 12...	4s
up-down.5me3t35c	up-down	Mar 17 2025, 12...	4s
up-down.5me3s5jt	up-down	Mar 17 2025, 12...	4s
up-down.5me3rocs	up-down	Mar 17 2025, 12...	4s
up-down.5me3qla	up-down	Mar 17 2025, 12...	4s
up-down.5me3q396	up-down	Mar 17 2025, 12...	4s

On the right, there is a 'Collect data' section with a 'Connect a device' button and a 'RAW DATA' section with a 'Click on a sample to load...' prompt.

The screenshot shows the 'Impulse design' page in the EDGE IMPULSE web interface. It features four main panels: 'Time series data', 'Spectral Analysis', 'Classification', and 'Output features'. The 'Time series data' panel has settings for input axes, window size, window increase, frequency, and zero-pad data. The 'Spectral Analysis' panel has a name field and a list of input axes. The 'Classification' panel has a name field, input features, and output features. The 'Output features' panel shows the selected output features and a 'Save Impulse' button.

Time series data

Input axes (9): accX, accY, accZ, gyrX, gyrY, gyrZ, magX, magY, magZ

Window size: 2,000 ms

Window increase (stride): 200 ms

Frequency (Hz): 100

Zero-pad data: ☒

Spectral Analysis

Name: Spectral features

Input axes (3): ☒ accX, ☒ accY, ☒ accZ, ☐ gyrX, ☐ gyrY, ☐ gyrZ, ☐ magX, ☐ magY, ☐ magZ

Classification

Name: Classifier

Input features: ☒ Spectral features

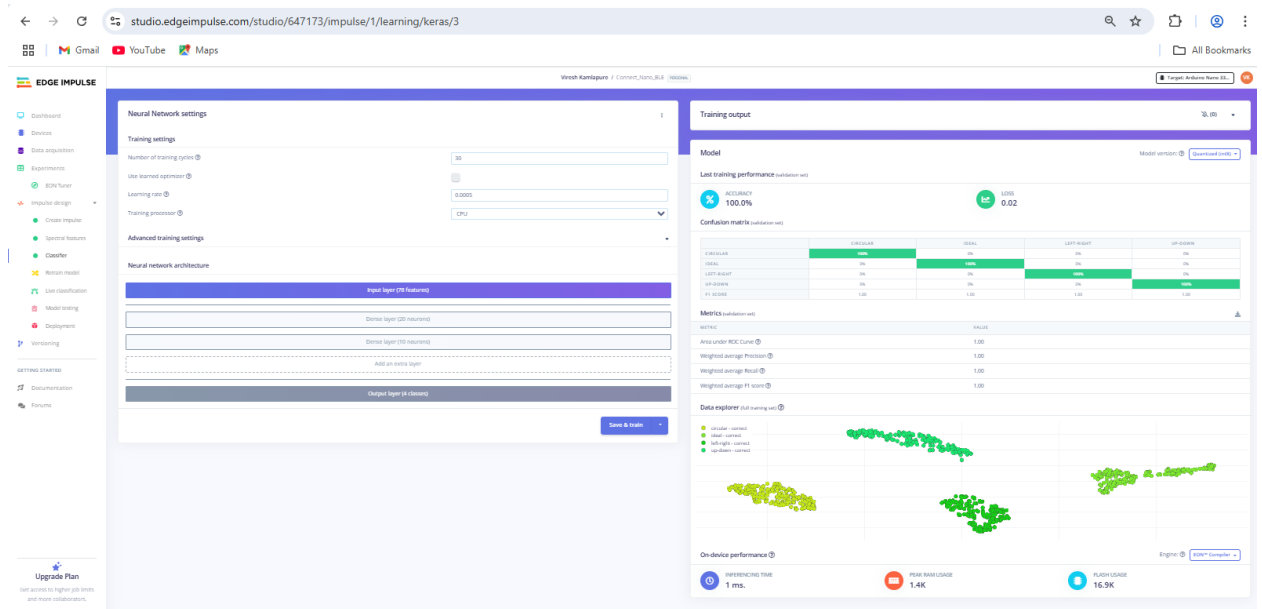
Output features: 4 (circular, ideal, left-right, up-down)

Output features

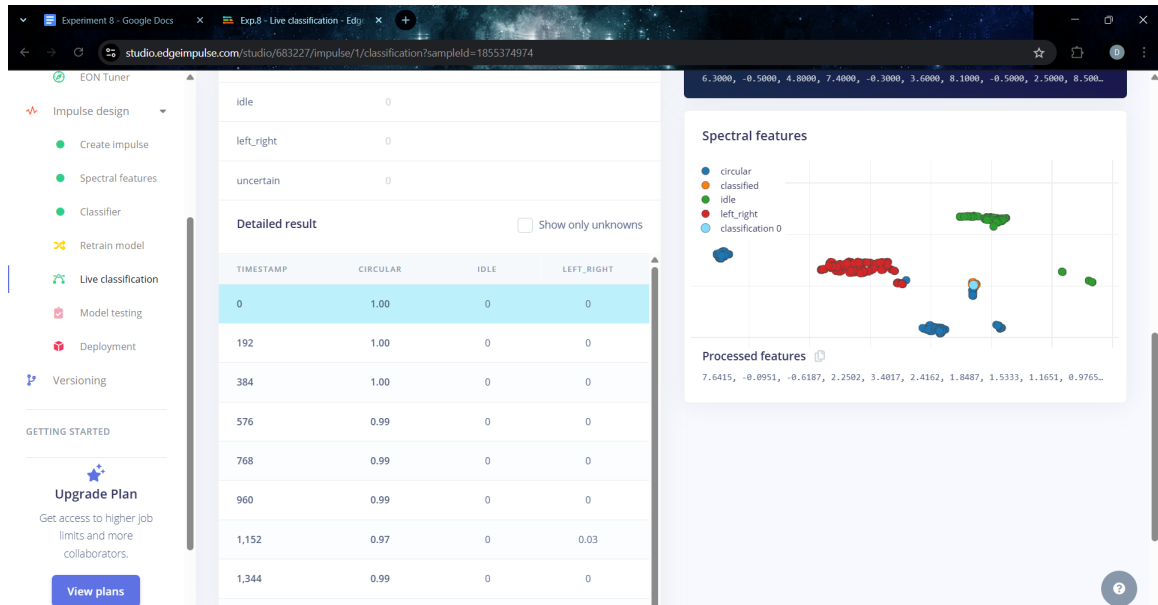
4 (circular, ideal, left-right, up-down)

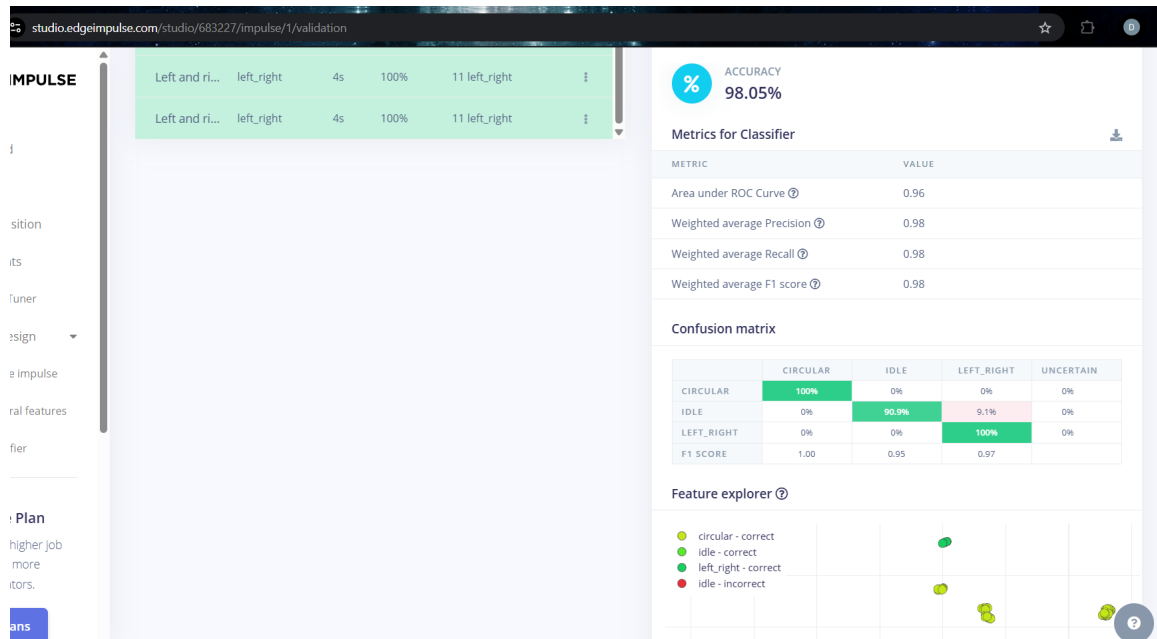
Save Impulse

3) Accuracy / Loss - Confusion Matrix – image



4) Validation Result – Image





5) Copy the code of Arduino Sketch

```

18 #include <nano_ble_project_inferencing.h>
19 #include <Arduino_LSM9DS1.h> //Click here to get the library: https://www.arduino.cc/reference/en/libraries/arduino\_lsm9ds1/
20
21 /* Constant defines ----- */
22 #define CONVERT_G_TO_MS2 9.80665f
23 /**
24  * When data is collected by the Edge Impulse Arduino Nano 33 BLE Sense
25  * firmware, it is limited to a 2G range. If the model was created with a
26  * different sample range, modify this constant to match the input values.
27  * See https://github.com/edgeimpulse/firmware-arduino-nano-33-ble-sense/blob/master/src/sensors/ei\_lsm9ds1.cpp
28  * for more information.
29  */
30 #define MAX_ACCEPTED_RANGE 2.0f
31
32 /*
33  ** NOTE: If you run into TFLite arena allocation issue.
34  **
35  ** This may be due to may dynamic memory fragmentation.
36  ** Try defining "-DEI_CLASSIFIER_ALLOCATION_STATIC" in boards.local.txt (create
37  ** if it doesn't exist) and copy this file to
38  ** "<ARDUINO_CORE_INSTALL_PATH>/arduino/hardware/<mbed_core>/<core_version>".
39  **
40  ** See
41  ** (https://support.arduino.cc/hc/en-us/articles/360012076960-where-are-the-installed-cores-located-)
42  ** to find where Arduino installs cores on your machine.
43  **
44  ** If the problem persists then there's not enough memory for this model and application.
45  */
46
47 /* Private variables ----- */
48 static bool debug_nn = false; // Set this to true to see e.g. features generated from the raw signal
49 static uint32_t run_inference_every_ms = 200;
50 static rtos::Thread inference_thread(osPriorityLow);
51 static float buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE] = { 0 };
52 static float inference_buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE];

```

```

53
54 /* Forward declaration */
55 void run_inference_background();
56
57 /**
58  * @brief Arduino setup function
59  */
60 void setup()
61 {
62     // put your setup code here, to run once:
63     Serial.begin(115200);
64     // comment out the below line to cancel the wait for USB connection (needed for native USB)
65     while (!Serial);
66     Serial.println("Edge Impulse Inferencing Demo");
67
68     if (!IMU.begin()) {
69         ei_printf("Failed to initialize IMU!\r\n");
70     }
71     else {
72         ei_printf("IMU initialized\r\n");
73     }
74
75     if (EI_CLASSIFIER_RAW_SAMPLES_PER_FRAME != 3) {
76         ei_printf("ERR: EI_CLASSIFIER_RAW_SAMPLES_PER_FRAME should be equal to 3 (the 3 sensor axes)\n");
77         return;
78     }
79
80     inference_thread.start(mbed::callback(&run_inference_background));
81 }
82
83 /**
84  * @brief Return the sign of the number
85  *
86  * @param number
87  * @return int 1 if positive (or 0) -1 if negative
88  */
89 float ei_get_sign(float number) {
90     return (number >= 0.0) ? 1.0 : -1.0;
91 }
92
93 /**
94  * @brief Run inferencing in the background.
95  */
96 void run_inference_background()
97 {
98     // wait until we have a full buffer
99     delay((EI_CLASSIFIER_INTERVAL_MS * EI_CLASSIFIER_RAW_SAMPLE_COUNT) + 100);
100
101     // This is a structure that smoothen the output result
102     // With the default settings 70% of readings should be the same before classifying.
103     ei_classifier_smooth_t smooth;
104     ei_classifier_smooth_init(&smooth, 10 /* no. of readings */, 7 /* min. readings the same */, 0.8 /* min. confidence */, 0.3 /* max anomaly */);
105
106     while (1) {
107         // copy the buffer
108         memcpy(inference_buffer, buffer, EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE * sizeof(float));
109
110         // Turn the raw buffer in a signal which we can the classify
111         signal_t signal;
112         int err = numpy::signal_from_buffer(inference_buffer, EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE, &signal);
113         if (err != 0) {
114             ei_printf("Failed to create signal from buffer (%d)\n", err);
115             return;
116         }
117
118         // Run the classifier
119         ei_impulse_result_t result = { 0 };
120
121         err = run_classifier(&signal, &result, debug_nn);
122         if (err != EI_IMPULSE_OK) {
123             ei_printf("ERR: Failed to run classifier (%d)\n", err);

```

```

123         ei_printf("ERR: Failed to run classifier (%d)\n", err);
124         return;
125     }
126
127     // print the predictions
128     ei_printf("Predictions ");
129     ei_printf("(DSP: %d ms., Classification: %d ms., Anomaly: %d ms.)",
130         result.timing.dsp, result.timing.classification, result.timing.anomaly);
131     ei_printf(": ");
132
133     // ei_classifier_smooth_update yields the predicted label
134     const char *prediction = ei_classifier_smooth_update(&smooth, &result);
135     ei_printf("%s ", prediction);
136     // print the cumulative results
137     ei_printf(" [ ");
138     for (size_t ix = 0; ix < smooth.count_size; ix++) {
139         ei_printf("%u", smooth.count[ix]);
140         if (ix != smooth.count_size + 1) {
141             ei_printf(", ");
142         }
143         else {
144             ei_printf(" ");
145         }
146     }
147     ei_printf("]\n");
148
149     delay(run_inference_every_ms);
150 }
151
152 ei_classifier_smooth_free(&smooth);
153 }
154
155 /**
156  * @brief      Get data and run inferencing
157  *
158  * @param[in]  debug  Get debug info if true
159  */
160
161 {
162     while (1) {
163         // Determine the next tick (and then sleep later)
164         uint64_t next_tick = micros() + (EI_CLASSIFIER_INTERVAL_MS * 1000);
165
166         // roll the buffer -3 points so we can overwrite the last one
167         numpy::roll(buffer, EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE, -3);
168
169         // read to the end of the buffer
170         IMU.readAcceleration(
171             buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3],
172             buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 2],
173             buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 1]
174         );
175
176         for (int i = 0; i < 3; i++) {
177             if (fabs(buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3 + i]) > MAX_ACCEPTED_RANGE) {
178                 buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3 + i] = ei_get_sign(buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3 + i]) * MAX_ACCEPTED_RANGE;
179             }
180         }
181
182         buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3] *= CONVERT_G_TO_MS2;
183         buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 2] *= CONVERT_G_TO_MS2;
184         buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 1] *= CONVERT_G_TO_MS2;
185
186         // and wait for next tick
187         uint64_t time_to_wait = next_tick - micros();
188         delay((int)floor((float)time_to_wait / 1000.0f));
189         delayMicroseconds(time_to_wait % 1000);
190     }
191 }
192
193 #if !defined(EI_CLASSIFIER_SENSOR) || EI_CLASSIFIER_SENSOR != EI_CLASSIFIER_SENSOR_ACCELEROMETER
194 #error "Invalid model for current sensor"
195 #endif
196

```

6) Screen shot of Arduino Terminal - Result


```
Starting inferencing in 2 seconds...
Sampling...
Predictions (DSP: 132.291000 ms., Classification: 0.580000 ms., Anomaly: 0ms.):
#Classification results:
  circular: 0.371094
  idle: 0.523437
  right_left: 0.042969
  up_down: 0.062500
Starting inferencing in 2 seconds...
Sampling...
Predictions (DSP: 133.824997 ms., Classification: 0.571000 ms., Anomaly: 0ms.):
#Classification results:
  circular: 0.000000
  idle: 0.996094
  right_left: 0.000000
  up_down: 0.000000
Starting inferencing in 2 seconds...
Sampling...
Predictions (DSP: 129.904007 ms., Classification: 0.571000 ms., Anomaly: 0ms.):
#Classification results:
  circular: 0.000000
  idle: 0.996094
  right_left: 0.000000
  up_down: 0.003906
```